

Q4. Distributed Concurrency Control

- Concurrency controlling techniques ensure that multiple transactions are executed simultaneously while maintaining the ACID properties of transactions and serializability in schedules.
- Concurrency control is provided in a database to →
 - enforce isolation among transactions
 - preserve database consistency
 - resolve read-write and write-read conflicts
- Locking based Concurrency control protocols use the concept of locking data items
- A lock is a variable associated with a data item that determines whether read/write operations can be performed on the data item.
- The operation's access request is decided based on the compatibility of ~~two~~ lock modes
- A read lock is compatible with another read lock
A write lock is not compatible with another read or write lock.
- Locking based concurrency control systems can either use one-phase or two-phase locking protocols.

One Phase Locking Protocol →

- In this method, each transaction locks an item before use and releases the lock as soon as it finishes using it.
- This locking method provides for maximum concurrency.
- But it does not always enforce serializability.

Two Phase Locking Protocol →

- In this method all the locking operations occur before the first lock release or unlock operation.

- The transaction comprises of two phases →

→ In the first phase, the transaction acquires all the locks it needs and does not release any lock.

This is called the expanding or growing phase.

→ In the second phase, the transaction releases the locks and cannot request any new locks.

This is called the shrinking phase.

Locks are usually maintained in a lock table.

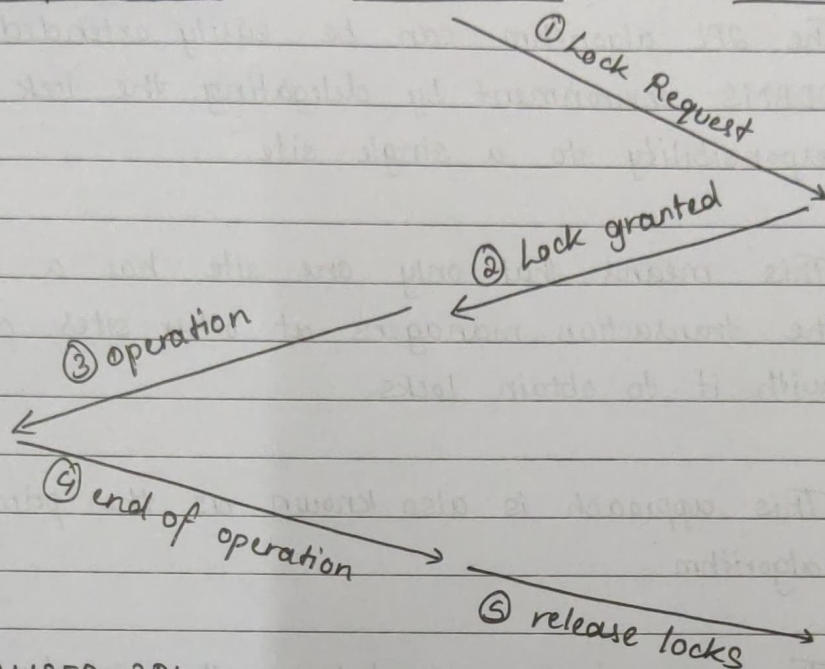
Centralized 2PL →

- The 2PL algorithm can be easily extended to the DDBMS environment by delegating the lock management responsibility to a single site.
- This means that only one site has a lock manager, the transaction managers at other sites communicate with it to obtain locks.
- This approach is also known as the primary site 2PL algorithm
- The communication is between the coordinating TM, the Lock manager at the central site and data processors at other participating sites.
- The participating sites are those ^{that store the data items} on which the operation(s) is to be carried out.
- The transaction manager (C2PL-TM) is written as an algorithm that runs forever and waits until a message arrives either from the application (with a transaction operation) or from the lock manager.
- The data processors (DP) and the lock manager (C2PL-LM) algorithms are written as procedures that are called when needed.

Data processors
at participating sites

Coordinating
TM

Participating
TM



CENTRALISED 2PL

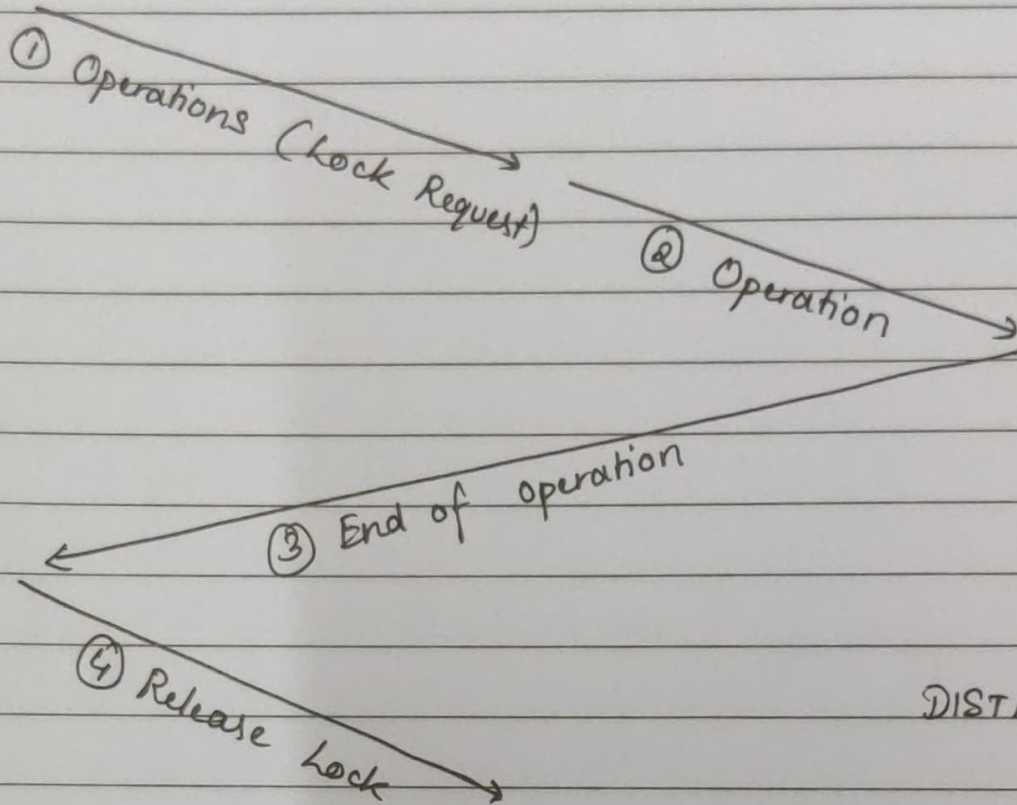
Distributed 2PL →

- D2PL requires the availability of lock manager at each site
- The D2PL-TM algorithm is similar to the C2PL-TM, with two major modifications
 - The messages sent to the central site lock manager in C2PL-TM are sent to all the sites participating sites in D2PL-TM.
 - Operations are not passed to the dataprocessors by the coordinating transaction manager, instead they are passed by the participating lock managers. So, the coordinating transaction manager does not wait for a lock request granted message

Coordinating
TM

Participating
Schedule^r

Participating
DM's



DISTRIBUTED

2PL